



Faculty of Engineering and Applied Science

SOFE 4790U Distributed Systems

CRN 44425

Lab #2

Alexander Campbell

100703650

Part 1:

The Gateway Routing pattern is a pattern where a single endpoint can be used to serve multiple services or multiple requests. This is done by placing a gateway in front of your application so that the connected client is only aware of the single endpoint occupied by the gateway. This is useful for three main reasons: exposing several versions of one service, having multiple instances of the same service, or when there's multiple services on one end point.

We can explore the usefulness by taking an e-commerce store as an example. Say a product owner would like to do some AB Testing for a new user interface, you may change the user interface on one version of the service to test if it increases sales by exposing the two versions of the service on one end point for your traffic to use. The gateway would then automatically distribute clients to each version of the website. Having multiple instances is also very useful for e-commerce as there are many regional differences you may want to consider such as high traffic times changing depending on the time zone. In this case, we would want to balance the load to ensure that servers located in high traffic areas have the appropriate amount of availability, while low traffic areas can be decreased and save on computational costs. The gateway would automatically redirect clients towards the endpoint that best suits their region. E-commerce stores also have many services such as search or checkout. Using this pattern, we can expose each end point so that the client is aware of each of the services available to them through the gateway.

The problems discussed in this can be boiled down to the fact that clients need to be easily redirected to different services without them having to have knowledge about the backend of the application. To utilize this pattern properly one must ensure that the gateway is correctly designed as otherwise it can lead to failure or a bottleneck of the entire system.

Part 2:

First the pods were initialized, then the web, experiment, and ambassador deployment yaml files were created, deployed, and exposed successfully

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ gcloud config set compute/zone northamerica-northeast1-b
Updated property [compute/zone].
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ gcloud container clusters create lab2 --num-nodes=3
Default change: VPC-native is the default mode during cluster creation for versions greater than 1.21.0-gke.1500. To create advanced routes based clusters, please pass the "--no-enable-ip-alias" flag
Default change: During creation of nodepools or autoscaling configuration changes for cluster versions greater than 1.24.1-gke.800 a default location policy is applied. For Spot and FVM it defaults to ANY, and for all other VM kinds a B
L4-NAT policy is used. To change the default values use the "--location-policy" flag.
Note: Your Pod address range ("--cluster-ip-v4-cidr") can accommodate at most 1008 node(s).
Creating cluster lab2 in northamerica-northeast1-b... Cluster is being health-checked (master is healthy)...done.
Created [https://containers.googleapis.com/v1/projects/dev-solstice-362019/zones/northamerica-northeast1-b/clusters/lab2].
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload/_gcloud/northamerica-northeast1-b/lab2?project=dev-solstice-362019
Kubernetes entry generated for lab2.
NAME: lab2
LOCATION: northamerica-northeast1-b
MASTER VERSION: 1.22.12-gke.2300
MASTER IP: 34.95.24.138
MACHINE TYPE: e2-medium
NODE VERSION: 1.22.12-gke.2300
NUM NODES: 3
STATUS: RUNNING
```

```
web-deployment.yaml x experiment-deployment.yaml nginx-ambassa... web-deployment.yaml x experiment-deployment.yaml x nginx-ambassa...
web-deployment.yaml > ...
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: web-deployment
5 spec:
6   replicas: 2
7   selector:
8     matchLabels:
9       run: web-deployment
10  template:
11    metadata:
12      labels:
13        run: web-deployment
14    spec:
15      containers:
16      - name: web-depoyment
17        image: mcr.microsoft.com/azuredocs/aci-helloworld
18        ports:
19        - containerPort: 80

experiment-deployment.yaml > ...
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: experiment-deployment
5 spec:
6   replicas: 2
7   selector:
8     matchLabels:
9       run: experiment-deployment
10  template:
11    metadata:
12      labels:
13        run: experiment-deployment
14    spec:
15      containers:
16      - name: experiment-depoyment
17        image: mcr.microsoft.com/azuredocs/aci-helloworld
18        ports:
19        - containerPort: 80
```

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl create -f web-deployment.yaml
deployment.apps/web-deployment created
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
web-deployment-6fdbb5c6bb-752fj     1/1     Running   0           10s
web-deployment-6fdbb5c6bb-nfd9f     1/1     Running   0           10s
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl expose deployment web-deployment --port=80 --type=ClusterIP --name web-deployment
service/web-deployment exposed
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl create -f experiment-deployment.yaml
deployment.apps/experiment-deployment created
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl expose deployment experiment-deployment --port=80 --type=ClusterIP --name experiment-deployment
service/experiment-deployment exposed
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl create configmap ambassador-config --from-file=conf.d
configmap/ambassador-config created
```

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl create -f ambassador-deployment.yaml
deployment.apps/ambassador-deployment created
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl expose deployment ambassador-deployment --port=80 --type=LoadBalancer
service/ambassador-deployment exposed
```


Then the pods, deployments, and services were checked to ensure that they were running without any errors

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get deployments
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
ambassador-deployment              2/2     2             2           22s
experiment-deployment              2/2     2             2           10m
web-deployment                     2/2     2             2           11m
```

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ambassador-deployment-66db4f7766-f9rzv  1/1     Running   0           30s
ambassador-deployment-66db4f7766-xthpz  1/1     Running   0           30s
experiment-deployment-7b47cbd668-181fq  1/1     Running   0           11m
experiment-deployment-7b47cbd668-nqdv9  1/1     Running   0           11m
web-deployment-6fdbb5c6bb-752fj         1/1     Running   0           12m
web-deployment-6fdbb5c6bb-nfd9f         1/1     Running   0           12m
```

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get services
NAME                                TYPE               CLUSTER-IP     EXTERNAL-IP   PORT(S)          AGE
ambassador-deployment              LoadBalancer      10.12.6.8      35.203.28.8   80:31372/TCP     66s
experiment-deployment              ClusterIP          10.12.7.229    <none>        80/TCP           11m
kubernetes                         ClusterIP          10.12.0.1      <none>        443/TCP          15m
web-deployment                     ClusterIP          10.12.7.194    <none>        80/TCP           12m
```

After the pods, deployments, and services were confirmed to be running correctly I opened up the ambassador-deployment IP address in my browser to find the website running and then ran the curl command many times and viewed the logs.



Welcome to Azure Container Instances!



```
a.bainhewcampbell@cloudshell:~ (dev-solstice-362019) $ curl http://35.203.28.8
<html>
<head>
<title>Welcome to Azure Container Instances</title>
</head>
<style>
hl {
  color: darkblue;
  font-family:arial, sans-serif;
  font-weight: lighter;
}
</style>
<body>

<div align="center">
<h1>Welcome to Azure Container Instances</h1>

<svg id="layer_1" data-name="Layer 1" xmlns="http://www.w3.org/2000/svg" viewBox="0 0 49.8 49.8" width="250px" height="250px">
<title>ContainerInstances_glb_H1</title>
<path d="M41.9,11.368A11.929,11.929,0,0,0,20.3,5.061a9.444,9.444,0,0,0-14.932,9.042,9.042,9.044,32H39.442A10.463,10.463,0,0,0,41.9,11.368Z" transform="translate(-0.1 -0.1)" fill="#ffff"/>
<path d="M41.9,11.368A11.929,11.929,0,0,0,20.3,5.061a9.444,9.444,0,0,0-14.932,9.042,9.042,9.044,32H39.442A10.463,10.463,0,0,0,41.9,11.368Z" transform="translate(-0.1 -0.1)" fill="#27a861" opacity="0.6" style="isolation:isolate"/>
</svg>
</div>
</body>
</html>
a.bainhewcampbell@cloudshell:~ (dev-solstice-362019) $ curl http://35.203.28.8
<html>
<head>
<title>Welcome to Azure Container Instances</title>
</head>
<style>
hl {
  color: darkblue;
  font-family:arial, sans-serif;
  font-weight: lighter;
}
</style>
<body>

a.bainhewcampbell@cloudshell:~ (dev-solstice-362019) $ kubectl logs -l run=web-deployment
:ffff:10.8.1.9 - - [05/Oct/2022:22:57:25 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:26 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.1.9 - - [05/Oct/2022:22:57:27 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.1.9 - - [05/Oct/2022:22:57:29 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:32 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:34 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:34 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:56:42 +0000] "GET / HTTP/1.0" 200 1663 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:56:42 +0000] "GET /favicon.ico HTTP/1.0" 404 150 "http://35.203.28.8/" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.1.9 - - [05/Oct/2022:22:57:22 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:25 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:26 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:33 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:34 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:35 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"

a.bainhewcampbell@cloudshell:~ (dev-solstice-362019) $ kubectl logs -l run=experiment-deployment
listening on port 80
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:29 +0000] "GET / HTTP/1.0" 200 1663 "-" "curl/7.74.0"
:ffff:10.8.0.5 - - [05/Oct/2022:22:57:35 +0000] "GET / HTTP/1.0" 304 - "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:105.0) Gecko/20100101 Firefox/105.0"
listening on port 80
```

Part 3:

Since the pods were set up from part 2, I created the yaml file, deployed, and exposed it. Then I verified that it was working correctly by running the specified commands.

```
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl create -f loadbalancer-deployment.yaml
deployment.apps/loadbalancer-deployment created
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get pods --output=wide
NAME                                READY    STATUS    RESTARTS   AGE    IP              NODE                                NOMINATED NODE    READINESS GATES
loadbalancer-deployment-6676f9ccf6-58w6x  1/1      Running   0           93s    10.8.2.8        gke-lab2-default-pool-2a786950-3rxq  <none>             <none>
loadbalancer-deployment-6676f9ccf6-g8xm6  1/1      Running   0           93s    10.8.0.8        gke-lab2-default-pool-2a786950-73bh  <none>             <none>
loadbalancer-deployment-6676f9ccf6-q4d8t  1/1      Running   0           93s    10.8.1.13       gke-lab2-default-pool-2a786950-pv4f  <none>             <none>
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl expose deployment loadbalancer-deployment --port=8080 --type=LoadBalancer
service/loadbalancer-deployment exposed
```

```
NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes                         ClusterIP           10.12.0.1     <none>         443/TCP          5m48s
loadbalancer-deployment            LoadBalancer       10.12.6.39    35.203.28.8    8080:32388/TCP   104s
^Ca_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ curl http://35.203.28.8:8080/dog
A quadruped of the genus Canis, esp. the domestic dog (C.familiaris).a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$
a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ curl http://35.203.28.8:8080/storey
See Story.a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$
```

Discussion:

As discussed in part 1, the major issues that the Gateway Routing pattern solve are when you want to expose multiple services on one end point, when there are multiple versions of the same service you want to expose to various users through one end point, and when you want multiple instances of a service on one endpoint. The solution that the Gateway Routing pattern provides is having a gateway service that will route the client to the various instances that an application may have.

The requirements of this pattern are that you have multiple deployments of some service(s) and a gateway service that will route the client to a deployment of their requested service(s).

In part 2 of the lab, this pattern was seen where we had the web deployment and the experimental deployment being kept behind the gateway of the ambassador deployment. The ambassador was used as a gateway such that you would be routed to one of the web or experimental deployments. This was seen after repeatedly using the curl command and viewing the logs of each of the deployments.

In part 3 of the lab, we deploy 3 images of the dictionary-server and use the readiness probe as a gateway. This gateway checks the readiness of all 3 deployed images every 5 seconds to ensure that they are healthy.

Design:

Autoscaling is a service that initiates new instances of your deployment automatically when your existing deployments are becoming overwhelmed or conversely closes deployments automatically when too many are idle or underutilized. This allows a platform to optimize the server utilization without having client performance degrade.

This differs from load balancing such that load balancers will balance the requests being taken in throughout existing deployments, while the autoscaler will create more deployments that the load balancer can route requests towards. Autoscaling is used when traffic to the platform is elastic as it can automatically scale up to meet the demands that a load balancer cannot do on its own, or scale down to save on computational costs.

How autoscaling is implemented can change depending on the business, however generally it involves the platform having a baseline amount of deployments that are scaled up

by a load balancing service when it finds there are not enough resources for the incoming requests.

Utilizing a kubernetes autoscaler guide, it was possible to set up an autoscaler and experiment with it and see the scaling done in real-time.

```
php-apache.yaml x
php-apache.yaml > {} spec > @ replicas
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: php-apache
5  spec:
6    selector:
7      matchLabels:
8        run: php-apache
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         run: php-apache
14     spec:
15       containers:
16       - name: php-apache
17         image: registry.k8s.io/hpa-example
18         ports:
19         - containerPort: 80
20       resources:
21         limits:
22           cpu: 500m
23         requests:
24           cpu: 200m
25   ---
26   apiVersion: v1
27   kind: Service
28   metadata:
29     name: php-apache
30     labels:
31       run: php-apache
32   spec:
33     ports:
34     - port: 80
35     selector:
36       run: php-apache

a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl apply -f https://k8s.io/examples/application/php-apache.yaml
deployment.apps/php-apache created
service/php-apache created
```

Configuration and deployment of yaml file

```
a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl run -i --tty load-generator --rm --image=busybox:1.28 --restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://php-apache; done"
If you don't see a command prompt, try pressing enter.
OK

a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get hpa
NAME          REFERENCE          TARGETS          MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache  <unknown>/1 (avg), <unknown>/60%  1        20       1          6m53s

a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get hpa php-apache --watch
NAME          REFERENCE          TARGETS          MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache  <unknown>/1 (avg), 0%/60%  1        20       1          7m27s
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       1          8m
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       4          8m15s
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       5          8m30s

^Ca_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5          2m49s
```

Using a command that generates load onto the deployments so that the autoscaler kicks in and creates more replicas to meet the need

```
a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl run -i --tty load-generator --rm --image=busybox:1.28 --restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://php-apache; done"
If you don't see a command prompt, try pressing enter.
OK

a_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get hpa php-apache --watch
NAME          REFERENCE          TARGETS          MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache  <unknown>/1 (avg), 0%/60%  1        20       1          7m27s
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       1          8m
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       4          8m15s
php-apache    Deployment/php-apache  <unknown>/1 (avg), 250%/60%  1        20       5          8m30s

^Ca_bainhecampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5          2m49s
```

```

a_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get hpa php-apache --watch
NAME          REFERENCE          TARGETS          MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache <unknown>/1 (avg), 43%/60% 1         20       5         9m54s
php-apache    Deployment/php-apache <unknown>/1 (avg), 0%/60% 1         20       5         10m
^Ca_bainhewcampbell@cloudshell:~ (dev-solstice-362019)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5          4m24s

```

After stopping the command we can see CPU Utilization go down