



Faculty of Engineering and Applied Science

SOFE 4790U Distributed Systems

CRN 44425

Lab # 3

Mamun Hossain - 100553073

Introduction:

This lab gives us an introduction to understand how to deploy a circuit breaker file and utilize the function as a service. We have also learned how to access FaaS services via the UI and invoke the requests accordingly.

Part 1: How to solve the problem and the requirements:

The problem at hand was the many factors that had affected cloud-hosted applications, and those factors were such as network latency, performance and availability of the underlying compute and storage systems and the network between them.

The way this problem was solved was by implementing health monitoring of the clusters and the nodes to each endpoint of the application. It will check if the application or services respond to the request of the health monitor and it will also analyze the endpoint of the tool or framework that performs the check.

The requirements to do this health monitor check is to implement an Agent that will have the appropriate certificates and in doing so, will send pings to the application endpoint to ensure that over the network everything is working, and will check for 200 (OK) status to be returned to ensure data errors are minimized.

Part 2: Procedure:

Here we can see that we initialize the setup with the required clusters and then clone the repository and input the dummy service containers and services and then push them with the docker command.

```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to velvety-rookery-362717.
Use 'veloud config set project [PROJECT ID]' to change to a different project.
mammutoclass@cloudshell:~ (velvety-rookery-362717)$ cd -
git clone https://github.com/GeorgeHoud3/SOFE47900-lab3.git
Cloning into 'SOFE47900-lab3'...
remote: Enumerating objects: 27, done.
remote: Counting objects: 100% (27/27), done.
remote: Compressing objects: 100% (20/20), done.
remote: Total 27 (delta 4), reused 24 (delta 4), pack-reused 0
Receiving objects: 100% (27/27), 5.77 KiB | 2.88 MiB/s, done.
Resolving deltas: 100% (4/4), done.
mammutoclass@cloudshell:~ (velvety-rookery-362717)$ cd ~/SOFE47900-lab3/part2/DummyServiceContainer
mammutoclass@cloudshell:~/SOFE47900-lab3/part2/DummyServiceContainer (velvety-rookery-362717)$ docker build -t us.gcr.io/~C
mammutoclass@cloudshell:~/SOFE47900-lab3/part2/DummyServiceContainer (velvety-rookery-362717)$ "C
Sending build context to Docker daemon 4.456kB
Step 1/7 : FROM node:carbon
carbon Pulling from library/node
146bd6a88618: Pull complete
9935d0c2ace: Pull complete
d0e7d6e4d6: Pull complete
e705a4c4fd31: Pull complete
cd7b722dbd4: Pull complete
445d0ee9214: Pull complete
db8fcd9db2fe: Pull complete
1d1510d1b3e: Pull complete
fd4993995f40: Pull complete
Digest: sha256:ia651bf74f05b8d0d3eb21a60der168a976108a287a74167ab593fc953aac34df
Status: Downloaded newer image for node:carbon
----> Seed4f3757f4
Step 2/7 : WORKDIR /usr/src/app
----> Running in 01a8d7522e7
Removing intermediate container 01a8d7522e7
----> 1f7a8c81290
Step 3/7 : COPY package*.json ./
----> 4b01f91749c6
Step 4/7 : RUN npm install
----> Running in 6db94b3382c0
npm notice created a lockfile as package-lock.json. You should commit this file.
npm WARN deprecated eslint@8.0.1: No repository field.
npm WARN dummy-service@1.0.1 No license field.

added 57 packages from 42 contributors and audited 57 packages in 1.986s
6 packages are looking for funding
run 'npm fund' for details

found 0 vulnerabilities
Removing intermediate container 6db94b3382c0
----> 5e14d7d4c0fd
Step 5/7 : COPY . .
----> 683b61de4f54
Step 6/7 : EXPOSE 80
----> Running in 760dfa4fb8c0
Removing intermediate container 760dfa4fb8c0
----> 159c27b08864
Step 7/7 : CMD [ "npm", "start" ]
----> Running in e54bc387a09
Removing intermediate container e54bc387a09
----> 346d4eae39
Successfully built 346d4eae39
Successfully tagged us.gcr.io/velvety-rookery-362717/dummy-service:latest
mammutoclass@cloudshell:~/SOFE47900-lab3/part2/DummyServiceContainer (velvety-rookery-362717)$
```

```
Successfully tagged us.gcr.io/velvety-rookery-362717/dummy-service:latest
mamunotuclass@cloudshell:~/SOFE4790U-lab3/part2/DummyServiceContainer (velvety-rookery-362717) $ gcloud projects list
PROJECT_ID: mamunhossain
NAME: MamunHossain
PROJECT_NUMBER: 418585279789

PROJECT_ID: sunny-effort-366715
NAME: SOFE4640 GoogleMapKey
PROJECT_NUMBER: 685349366104
```

```
PROJECT_ID: velvety-rookery-362717
NAME: My First Project
PROJECT_NUMBER: 38427508957
mammuniculclass@cloudshell:~/-/SOFE4790U-lab3/part2/DummyServiceContainer (velvety-rookery-362717)$ docker push us.gcr.io/velvety-rookery-362717/dummyservice
Using default tag: latest
The push refers to repository [us.gcr.io/velvety-rookery-362717/dummyservice]
fdb15b87583c: Pushed
abd302d5b471: Pushed
5f873b52387a: Pushed
17e0dfafdf13: Pushed
423451ed44f2: Layer already exists
b2aa8f5d6633: Layer already exists
88601a85cc11: Layer already exists
42f9c2f3c08e: Layer already exists
99e8bd3efaaaf: Layer already exists
bee1e39d7c3a: Layer already exists
1f59a4b2e206: Layer already exists
0ca7f54856c0: Layer already exists
ebb9ae013834: Layer already exists
latest: digest: sha256:a70f5b4579e2eb15b7ddeb145f33015a63db97bfb34533bfbf770d151aa38e6 size: 3048
mammuniculclass@cloudshell:~/-/SOFE4790U-lab3/part2/DummyServiceContainer (velvety-rookery-362717)$
```

[illegible]

```

12   run: backup-deployment
13
14 metadata:
15   labels:
16     run: backup-deployment
17
18 spec:
19   containers:
20     - name: backup-deployment
21       image: us.gcr.io/vast-alcove-367116/dummyservice
22       ports:
23         - containerPort: 80
24       livenessProbe:
25         httpGet:
26           # The /alive endpoint is the one we will not touch in our test case, a
27           path: /alive
28           port: 80

```

Here, we can see after we create the configmap, we can use the curl command to execute the circuit breaker UP and check it in the local machine in the last line of each response.

```
mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -v http://34.152.45.250
* Trying 34.152.45.250:80...
* Connected to 34.152.45.250 (34.152.45.250) port 80 (#0)
> GET / HTTP/1.1
> Host: 34.152.45.250
> User-Agent: curl/7.74.0
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.13.7
< Date: Sat, 29 Oct 2022 23:56:54 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 28
< Connection: keep-alive
< X-Powered-By: Express
< ETag: W/"1c-GjMsSn4PdftbTWqSpIr4BB7pY"
<
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -v http://34.152.45.250
* Trying 34.152.45.250:80...
* Connected to 34.152.45.250 (34.152.45.250) port 80 (#0)
> GET / HTTP/1.1
> Host: 34.152.45.250
> User-Agent: curl/7.74.0
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.13.7
< Date: Sat, 29 Oct 2022 23:57:45 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 28
< Connection: keep-alive
< X-Powered-By: Express
< ETag: W/"1c-GjMsSn4PdftbTWqSpIr4BB7pY"
<
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -v http://34.152.45.250
* Trying 34.152.45.250:80...
* Connected to 34.152.45.250 (34.152.45.250) port 80 (#0)
> GET / HTTP/1.1
> Host: 34.152.45.250
> User-Agent: curl/7.74.0
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.13.7
< Date: Sat, 29 Oct 2022 23:57:47 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 28
< Connection: keep-alive
< X-Powered-By: Express
< ETag: W/"1c-GjMsSn4PdftbTWqSpIr4BB7pY"
<
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$
```

While we see the last line remaining unchanged after each test, but the dummy service having a different line in each response, check test the circuit breaker one more time.

```
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -d "" -s -D - http://34.152.54.18/fakeerrormodeon
HTTP/1.1 200 OK
X-Powered-By: Express
Content-Type: text/html; charset=utf-8
Content-Length: 18
ETag: W/"12-aAbJlNj0ezsrBwtXVaCSzPVQhlc"
Date: Sat, 29 Oct 2022 23:58:34 GMT
Connection: keep-alive

OK FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -v http://34.152.45.250
* Trying 34.152.45.250:80...
* Connected to 34.152.45.250 (34.152.45.250) port 80 (#0)
> GET / HTTP/1.1
> Host: 34.152.45.250
> User-Agent: curl/7.74.0
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.13.7
< Date: Sat, 29 Oct 2022 23:58:41 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 28
< Connection: keep-alive
< X-Powered-By: Express
< ETag: W/"1c-GjMsSn4PdftbTWqSpIr4BB7pY"
<
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -d "" -s -D - http://34.152.54.18/fakeerrormodeon
HTTP/1.1 200 OK
X-Powered-By: Express
Content-Type: text/html; charset=utf-8
Content-Length: 18
ETag: W/"12-aAbJlNj0ezsrBwtXVaCSzPVQhlc"
Date: Sun, 30 Oct 2022 00:00:14 GMT
Connection: keep-alive

OK FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$ curl -v http://34.152.45.250
* Trying 34.152.45.250:80...
* Connected to 34.152.45.250 (34.152.45.250) port 80 (#0)
> GET / HTTP/1.1
> Host: 34.152.45.250
> User-Agent: curl/7.74.0
> Accept: */*
>
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< Server: nginx/1.13.7
< Date: Sun, 30 Oct 2022 00:00:16 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 28
< Connection: keep-alive
< X-Powered-By: Express
< ETag: W/"1c-GjMsSn4PdftbTWqSpIr4BB7pY"
<
* Connection #0 to host 34.152.45.250 left intact
SOMERESPONSE FROM 10.76.0.14mammutoclass@cloudshell:~/SOFE4790U-lab3/part2 (velvety-rookery-362717)$
```

Part 3:

We initialized the cluster admin role binding to then deploy OpenFaaS through the GKE, afterwards installing FaaS and as we can see below the confirmation of that and Arkade. We then use the commands to test that FaaS has been successfully initialized.

```

Welcome to Cloud Shell. Type "help" to get started.
Your Cloud Platform project in this session is set to velvety-rookery-362717.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
mmamoutclass@cloudshell: (velvety-rookery-362717) $ kubectl create clusterrolebinding "cluster-admin-${whoami}" \
error: required flag(s) "clusterrole" not set
mmamoutclass@cloudshell: (velvety-rookery-362717) $ kubectl create clusterrolebinding "cluster-admin-${whoami}" \
--clusterrole=cluster-admin \
--user=${gcloud config get-value core/account}
Your active configuration is: [cloudshell-32652]
error: exactly one role is required, got 3
See "kubectl create clusterrolebinding -h" for help and examples
mmamoutclass@cloudshell: (velvety-rookery-362717) $ kubectl config use-context kubernetess-admin@kubernetess
error: context exists with the name "kubernetess-admin@kubernetess"
mmamoutclass@cloudshell: (velvety-rookery-362717) $ kubectl create clusterrolebinding "cluster-admin-${whoami}" \
--clusterrole=cluster-admin \
--user=${gcloud config get-value core/account}
Your active configuration is: [cloudshell-32652]
error: exactly one role is required, got 3
See "kubectl create clusterrolebinding -h" for help and examples
mmamoutclass@cloudshell: (velvety-rookery-362717) $ kubectl create clusterrolebinding "cluster-admin-${whoami}" \
--clusterrole=cluster-admin \
--user=${gcloud config get-value core/account}
Your active configuration is: [cloudshell-32652]
clusterrolebinding.rbac.authorization.k8s.io/cluster-admin-mmamoutclass created
mmamoutclass@cloudshell: (velvety-rookery-362717) $ curl -sSlf https://gh.get-arkade.dev/ | sudo sh
x86_64
Downloading package https://github.com/alexellis/arkade/releases/download/0.8.50/arkade as /tmp/arkade
Download complete.

Running with sufficient permissions to attempt to move arkade to /usr/local/bin
New version of arkade installed to /usr/local/bin
Creating alias 'ark' for 'arkade'.

```

The installation process of FaaS and ArKade.

[illegible]

Successful verification of FaaS.

```

CLI:
commit: 8820d8e4a15dab900d8a7e8fc271851ccb94012e
version: 0.14.11
mamunotuclass@cloudshell:~ (velvety-rookery-362717)$
mamunotuclass@cloudshell:~ (velvety-rookery-362717)$ kubectl -n openfaas get deployments -l "release=openfaas, app=openfaas"
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
alertmanager  1/1      1              1            15s
basic-auth-plugin  1/1      1              1            15s
gateway       1/1      1              1            15s
nats          1/1      1              1            15s
prometheus    1/1      1              1            15s
queue-worker   1/1      1              1            15s
mamunotuclass@cloudshell:~ (velvety-rookery-362717)$

```

```
0x1f
arguments: 8d37df6d11fdcb0504017e8d9331821cbb44013e
```

Downloaded from <http://ajphaphysocpharm.sagepub.com/> at 11:01 11 September 2014

Here we are building the main.yml file after editing the FaaS main/handler.js file.

```
Digest: sha256:d4b15b3d48f42059a15bd659be60afe21762aae9d6cbea6f124440895c27db68
Status: Downloaded newer image for node:12-alpine
--> bb6d28039b8c
Step 3/31 : ARG TARGETPLATFORM
--> Running in 879e8a264f37
Removing intermediate container 879e8a264f37
--> d1d071a9a8bf
Step 4/31 : ARG BUILDPLATFORM
--> Running in 003bf19040f4
Removing intermediate container 003bf19040f4
--> be0058da2ff2
Step 5/31 : COPY --from=watchdog /fwatcdog /usr/bin/fwatcdog
--> a445e03b75ca
Step 6/31 : RUN chmod +x /usr/bin/fwatcdog
--> Running in 631a183c9e5b
Removing intermediate container 631a183c9e5b
--> dcb64a05fe38
Step 7/31 : RUN apk --no-cache add curl ca-certificates && addgroup -S app && adduser -S -g app app
--> Running in 81e4b0a9eef0
fetch https://dl-cdn.alpinelinux.org/alpine/v3.15/main/x86_64/APKINDEX.tar.gz
fetch https://dl-cdn.alpinelinux.org/alpine/v3.15/community/x86_64/APKINDEX.tar.gz
(1/5) Installing ca-certificates (20220614-r0)
(2/5) Installing brotli-libs (1.0.9-r5)
(3/5) Installing nghttp2-libs (1.46.0-r0)
(4/5) Installing libcurl (7.80.0-r4)
(5/5) Installing curl (7.80.0-r4)
Executing busybox-1.34.1-r5.trigger
```

Edited OpenFaaS handler.js file.

```
OpenFaaS > main >  handler.js > ...
1 'use strict'
2
3 module.exports = async (event, context) => {
4   var parameters=JSON.stringify(event.body)
5   return context
6   .status(200)
7   .succeed(parameters)
8 }
```

We then run the command and get the string input as desired

```
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ curl http://34.95.52.188:8080/function/main -H 'Content-Type: application/json' -d '{ "Name": "Square", "Color": "Red", "Dimensions": 2 }'
{"Name":"Square","Color":"Red","Dimensions":2}mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$
```

After repeating the same processes as the previous task, we defined the decorator with the file provided, making minor alterations with the IP and the file itself to make it work and give the desired output shown below.

```
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ cd ~/OpenFaaS
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ faas-cli new --lang node12 --prefix us.gcr.io/vaast-alcove-367116 decorator
Folder decorator already exists
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ faas-cli deploy -f decorator.yml
Deploying decorator
WARNING! You are not using an encrypted connection to the gateway, consider using HTTPS.
Deployed. 202 Accepted.
URL: http://34.95.52.188:8080/function/decorator
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ docker push us.gcr.io/vaast-alcove-367116/decorator
Using default tag: latest
The push refers to repository [us.gcr.io/vaast-alcove-367116/decorator]
60f0ed20477: Layer already exists
3480350834f: Layer already exists
4420ba9d85e: Layer already exists
b1e158b37e8: Layer already exists
16f70ee128e: Layer already exists
60b1fde3ab2: Layer already exists
182f2d2123f: Layer already exists
95f3cadd126: Layer already exists
688ba2554c1: Layer already exists
27ba3a5e9d5: Layer already exists
b0cc70513b0a: Layer already exists
7f70ca34399: Layer already exists
f6810f5920c: Layer already exists
d780c46c602: Layer already exists
4fc24d88285: Layer already exists
latest digest: sha256:c02736a3e12d74409ce23bf6c0d8f80148f75427b2d0d4ac13a7a2cd85c3 also: 3663
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ faas-cli deploy -f decorator.yml
Deploying decorator
WARNING! You are not using an encrypted connection to the gateway, consider using HTTPS.
Deployed. 202 Accepted.
URL: http://34.95.52.188:8080/function/decorator
mamunotucias@cloudshell:~/OpenFaaS (vaast-alcove-367116)$ curl http://34.95.52.188:8080/function/decorator -H 'Content-Type: application/json' -d '{ "Name": "Square", "Dimensions": 2 }'
{"Name":"Square","Dimensions":2}
$ bash: cd: command not found
```

File alteration of decorator/handler.js.

```
OpenFaaS > decorator > handler.js > <unknown> > ...
1  'use strict'
2  const request = require('sync-request');
3
4  module.exports = async (event, context) => {
5      var obj = event.body;
6      if (obj['Name'] === undefined) {
7          obj['Name'] = 'Nameless';
8      }
9      if (obj['Color'] === undefined) {
10         obj['Color'] = 'Transparent';
11     }
12     var res = request('POST', 'http://34.95.52.188:8080/function/decorator', {
13         body: JSON.stringify(obj)
14     });
15     console.log(res["body"].toString())
16     return context.status(200).succeed(res["body"].toString('utf8', 1, res["body"].length-
17 1).replace(/\\\\"/g, '\"'));
18 }
```

Discussion:

In summary, the main problem at hand was to understand how latency works and how the applications need to have health monitoring checks to ensure that it is working to its highest functionality. We need to understand that we cannot always monitor the hardware in the system for every device, as that can end up being costly and wasting time in the long run. With the health monitoring system we had discussed in part 1, we can use it effectively to ensure that we can understand and monitor the network latency, performance and availability of the underlying compute and storage systems with the networks between it.

We will be gathering different values and monitor the health checks as we did in the lab. In summary for part 1, we were able to confirm the need for implementing health monitoring of clusters and the nodes to each endpoint of the application to ensure that all the checkpoints were in place to monitor the node health. The way we were able to accomplish this with part 2 and 3 was by implementing several services: Backup service, implemented service, and a circuit breaker. We had used the circuit breaker several times and then paired it with a dummy service to ensure that the circuit breaker was working as it was intended to do. The way this would work is that the circuit breaker would be tested to make sure that the endpoints have the same value so it can be deemed to be a healthy endpoint.

For part 3, we had used the decorator function and the main function to distinguish the difference between the two and their interactions with the application and to view the values that get sent to the applications via JSON files that get converted into strings. The decorator endpoint ensures the validity of the checkpoint for the health monitoring by responding to any errors in the system and fixing them accordingly.

Design:

As we know, Kubernetes provides persistent volumes and it is necessary to have these volumes in a distributed system because if these volumes were not present, any data that has been previously created or altered will become null and void when the container restarts its processes. An example of when persistent volumes are needed is when let's say for example you want to be a developer at an enterprise, you want to make sure that your project is scalable. This way, if the consumer requests for something that can have several new environments to ensure that the product they want can run in different environments to their hearts desire. As we can see below, we will show how to configure a YAML file to implement the example and then running it and test it with persistent volume. We will be using "minikube" based off a guide we had found on the Internet doing our own research.

We install minikube using the commands and then check for all the pods to be working to make sure that the application will be working.

```
mamunotucass@cloudshell:~ (vast-alcove-367116)$ minikube start
* minikube v1.27.0 on Debian 11.5 (amd64)
- MINIKUBE_FORCE_SYSTEMD=true
- MINIKUBE_HOME=/google/minikube
- MINIKUBE_WANTUPDATENOTIFICATION=false
! Kubernetes 1.25.0 has a known issue with resolv.conf. minikube is using a workaround that should work for most use cases.
! For more information, see: https://github.com/kubernetes/kubernetes/issues/112135
* Using the docker driver based on existing profile
* Starting control plane node minikube in cluster minikube
* Pulling base image ...
* Updating the running docker "minikube" container ...
* Preparing Kubernetes v1.25.0 on Docker 20.10.17 ...
- kubelet.cgroups-per-qos=false
- kubelet.enforce-node-allocatable=""
* Verifying Kubernetes components...
- Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl get pods
No resources found in default namespace.
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl get pods -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  coredns-565d847f94-n6fwd              1/1     Running   1 (57s ago)  3m36s
kube-system  etcd-minikube                          1/1     Running   2 (57s ago)  3m48s
kube-system  kube-apiserver-minikube               1/1     Running   2 (47s ago)  3m48s
kube-system  kube-controller-manager-minikube      1/1     Running   2 (57s ago)  3m48s
kube-system  kube-proxy-v9wj7                      1/1     Running   2 (57s ago)  3m37s
kube-system  kube-scheduler-minikube               1/1     Running   1 (62s ago)  3m48s
kube-system  storage-provisioner                   1/1     Running   3 (43s ago)  3m46s
mamunotucass@cloudshell:~ (vast-alcove-367116)$ minikube dashboard
* Enabling dashboard ...
- Using image docker.io/kubernetes/metrics-scraper:v1.0.8
- Using image docker.io/kubernetes/dashboard:v2.6.0
* Verifying dashboard health ...
* Launching proxy ...
* Verifying proxy health ...
* Opening http://127.0.0.1:41731/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:proxy/ in your default browser...
http://127.0.0.1:41731/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:proxy/
```

We can now enter the node and create persistent volumes as seen below that the pods can utilize as persistent volumes.

```
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl apply -f https://k8s.io/examples/pods/storage/pv-volume.yaml
persistentvolume/task-pv-volume created
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl apply -f https://k8s.io/examples/pods/storage/pv-claim.yaml
persistentvolumeclaim/task-pv-claim created
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl apply -f https://k8s.io/examples/pods/storage/pv-pv-pod
error: unable to read URL "https://k8s.io/examples/pods/storage/pv-pv-pod", server reported 404 Not Found, status code=404
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl apply -f https://k8s.io/examples/pods/storage/pv-pod.yaml
pod/task-pv-pod created
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl get pvc task-pv-claim
NAME          STATUS    VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   AGE
task-pv-claim Bound     task-pv-volume   10Gi       RWO            manual         45s
mamunotucass@cloudshell:~ (vast-alcove-367116)$
```

From here, we can see the node being re entered by the command as seen below to ensure the persistent volume is being used as needed as shown created earlier into the node.

```
NAME          STATUS    VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   AGE
task-pv-claim Bound     task-pv-volume   10Gi       RWO            manual         45s
mamunotucass@cloudshell:~ (vast-alcove-367116)$ kubectl exec -it task-pv-pod -- /bin/bash
root@task-pv-pod:/# apt update
Get:1 http://deb.debian.org/debian bullseye InRelease [116 kB]
Get:2 http://deb.debian.org/debian-security bullseye-security InRelease [48.4 kB]
Get:3 http://deb.debian.org/debian bullseye-updates InRelease [44.1 kB]
Get:4 http://deb.debian.org/debian bullseye/main amd64 Packages [8184 kB]
Get:5 http://deb.debian.org/debian-security bullseye-security/main amd64 Packages [193 kB]
Get:6 http://deb.debian.org/debian bullseye-updates/main amd64 Packages [14.6 kB]
Fetched 8600 kB in 2s (4735 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
1 package can be upgraded. Run 'apt list --upgradable' to see it.
root@task-pv-pod:/#
```

Below are the yaml files in the editor

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 3Gi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: task-pv-pod
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: task-pv-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: task-pv-storage
```